

MedDialogue: General-Purpose Medical Dialogue Fine-Tuning Framework

Project Information

Lead/Mentor: Frederick Gyasi

Contributors: Frederick Gyasi

Institution: Medical University of South Carolina, Biomedical Informatics Center, Clinical NLP Lab

Current Funding: N/A

Future Funding: N/A

IRB #: N/A

RMID: N/A

SPARCRequest: N/A

Project Summary

MedDialogue is a general-purpose framework for fine-tuning large language models (LLMs) on ANY medical dialogue task through conversational learning. The framework is task-agnostic and contains no domain-specific code - users define their own tasks, output fields, and question templates.

Unlike traditional instruction fine-tuning that produces brittle models failing on question variations, MedDialogue trains models through natural multi-turn clinical dialogues, enabling robust handling of varied questions, follow-ups, and context.

Key Research Questions

1. **Can conversational fine-tuning produce more robust medical AI models than traditional instruction fine-tuning?**

2. How effective are multi-turn dialogue patterns in teaching clinical reasoning compared to single-turn approaches?
-

Framework Capabilities

Core Architecture

Component	Description
Task Definition	User-defined via <code>TaskConfig</code> (task name, input field, output fields, question templates)
Conversation Generation	Configurable via <code>ConversationConfig</code> (turn ratios, typos, validation split)
Question Combination	16 styles: 7 grammatical variations + 9 logical reasoning patterns
Output Formats	TEXT, JSON, XML, Markdown with configurable weighted ratios
Training Method	LoRA/QLoRA fine-tuning with Unslot optimization
Data Mapping	1:1 mapping (each clinical note generates ONE training conversation)

Supported Models

Model Family	Variants	Max Sequence Length
LLaMA	3.1, 3.2 (8B, 70B)	32,768 tokens
Phi	Phi-4 (14B)	16,384 tokens
Mistral	7B, Nemo	32,768 tokens
Qwen	2.5 (7B, 14B, 32B)	32,768 tokens

Question Generation Styles

Category	Styles	Description
Grammatical (1-7)	Sequential, Conjunctions, Transitional, Numbered, Natural, Directive, Mixed	Work with any question order
Logical (8-16)	Causal, Conditional, Prioritized, Comparative, Temporal, Hierarchical, Dialectical, Exploratory, Dependency	Require ordered questions for reasoning flow

Example Use Case: Pediatric Malnutrition Assessment

The repository includes a demonstration using pediatric malnutrition assessment (ASPEN/WHO/CDC guidelines) to show how to use MedDialogue:

File	Purpose
<code>train_malnutrition.py</code>	Training script demonstrating TaskConfig setup
<code>evaluate_malnutrition.py</code>	Evaluation script with classification metrics
<code>malnutrition_system_prompts.py</code>	Domain-specific system prompts

This example demonstrates:

- 11 clinical output fields
- 110 question templates (10 per field)
- Field ordering following clinical logic (assess, diagnose, plan)
- Multi-turn clinical dialogue patterns

Resources

Code Repository

GitHub: <https://github.com/gyasifred/meddialogue>

Core Framework Modules

Module	Function
	Main <code>MedDialogue</code> class - high-level training interface

Module	Function
<code>meddialogue/core.py</code>	
<code>meddialogue/config.py</code>	Configuration dataclasses (<code>TaskConfig</code> , <code>ConversationConfig</code> , <code>TrainingConfig</code> , <code>LoRAConfig</code> , <code>ModelConfig</code>)
<code>meddialogue/data_prep.py</code>	<code>DataPrep</code> class - conversation generation with 16 question styles
<code>meddialogue/train.py</code>	<code>Trainer</code> class - training pipeline with Unsloth
<code>meddialogue/inference.py</code>	<code>InferencePipeline</code> - single/multi-turn inference
<code>meddialogue/models.py</code>	<code>ModelRegistry</code> - model loading and LoRA application
<code>meddialogue/utils.py</code>	Text processing, output formatting, parsing utilities
<code>meddialogue/safety.py</code>	Safety stubs for backward compatibility

Data Requirements

Users provide CSV data with:

Required	Description
Input field	Clinical text column (configurable, default: <code>clinical_note</code>)
Output fields	User-defined columns matching <code>TaskConfig.output_fields</code>

Acknowledgements

Lead/Mentor: Frederick Gyasi (gyasi@musc.edu)

Contributors: Frederick Gyasi

Past Contributors: N/A

Institution: Medical University of South Carolina, Biomedical Informatics Center,
Clinical NLP Lab

Current Funding: N/A

Future Funding: N/A

License: MIT

Technical Dependencies

Package	Purpose
Unsloth	Fast LLM fine-tuning with memory optimization
Transformers	Hugging Face model infrastructure
PEFT	Parameter-efficient fine-tuning (LoRA)
PyTorch	Deep learning framework
Datasets	HuggingFace dataset handling
pandas	Data processing
scikit-learn	Evaluation metrics

Contact: gyasi@musc.edu

Medical University of South Carolina, 2025